

▼ Lab2A - Introduction to PyTorch

Tensor (`torch.tensor`) is the data structure used in PyTorch to build a deep learning system. Tensors are similar to NumPy's `ndarrays` , with the addition being that Tensors can also be used on a GPU to accelerate computing.

Objectives:

In this lab, you learn how to

- Create tensors in PyTorch
- Perform mathematical operation on tensors
- Convert between PyTorch tensor and Numpy array
- Reshape a PyTorch tensor
- Transfer tensor to and from GPU

Table of Content:

1. [Creating tensors](#)
2. [Tensor operations](#)
3. [Indexing](#)
4. [Reshaping tensors](#)
5. [CUDA Tensors](#)
6. [Exercise](#)

Reference:

- [PyTorch Official Tutorial: What is PyTorch](#)
-

▼ 1. Creating tensors

```
1 import torch
```

Construct a 5x3 matrix, uninitialized

```
1 x = torch.empty(5, 3)
2 print(x)

tensor([[1.5308e+21, 3.0823e-41, 3.3631e-44],
        [0.0000e+00, nan, 3.0823e-41],
        [1.1578e+27, 1.1362e+30, 7.1547e+22],
        [4.5828e+30, 1.2121e+04, 7.1846e+22],
        [9.2198e-39, 7.0374e+22, 2.8865e+20]])
```

Construct a tensor filled with random numbers from a uniform distribution on the interval $[0, 1)$.

```
1 x = torch.rand(5, 3)
2 print(x)

tensor([[0.7521, 0.2607, 0.5136],
        [0.0039, 0.1750, 0.5553],
        [0.7461, 0.3617, 0.8170],
        [0.6385, 0.8919, 0.1916],
        [0.4943, 0.4376, 0.7084]])
```

Construct a tensor filled with random numbers from a normal distribution with mean 0 and variance 1.

```
1 x = torch.randn(5, 3)
2 print(x)

tensor([[ -0.6189,  0.4508,  0.8886],
        [-1.0557,  1.2414, -0.3561],
        [ 0.9360, -1.5024,  0.9378],
        [ 0.2679, -0.2584, -0.0786],
        [ 0.2679, -0.2584, -0.0786]])
```

```
[-1.6949, -0.3799,  0.4828]])
```

Construct a matrix filled with zeros and of dtype long

```
1 x = torch.zeros(5, 3, dtype=torch.long)
2 print(x)
```

```
tensor([[0, 0, 0],
        [0, 0, 0],
        [0, 0, 0],
        [0, 0, 0],
        [0, 0, 0]])
```

```
1 x = torch.ones(5, 3)
2 print(x)
```

```
tensor([[1., 1., 1.],
        [1., 1., 1.],
        [1., 1., 1.],
        [1., 1., 1.],
        [1., 1., 1.]])
```

▼ 2. Tensor Operations

Size of tensors

```
1 x = torch.rand(5, 3)
2 print(x)
```

```
3 print(x.size()) # torch.Size is actually a tuple
4 print(x.shape)
```

```
tensor([[0.5162, 0.1062, 0.2726],
```

```

tensor([[0.3403, 0.4003, 0.2720],
        [0.6868, 0.6290, 0.0590],
        [0.1826, 0.4690, 0.7819],
        [0.8968, 0.0451, 0.6952],
        [0.8669, 0.0513, 0.2640]])
torch.Size([5, 3])
torch.Size([5, 3])

```

▼ Addition

There are multiple syntaxes for operations.

Addition: syntax 1

```

1 x = torch.rand(3, 2)
2 print('x:\n', x)
3 y = torch.rand(3, 2)
4 print('y:\n', y)
5
6 z = x + y
7 print('x+y:\n', z)

```

```

x:
tensor([[0.3215, 0.2230],
        [0.2855, 0.1580],
        [0.3406, 0.8168]])
y:
tensor([[0.8634, 0.2672],
        [0.2046, 0.7790],
        [0.7036, 0.6863]])
x+y:
tensor([[1.1849, 0.4902],
        [0.4901, 0.9369],
        [1.0442, 1.5032]])

```

Addition: syntax 2

```
1 z = torch.add(x, y)
2 print('x+y:\n', z)
```

```
x+y:
tensor([[1.1849, 0.4902],
        [0.4901, 0.9369],
        [1.0442, 1.5032]])
```

Addition: syntax 3 (in-place)

- Any operation that mutates a tensor in-place is post-fixed with an `_`. For example: `x.copy_(y)`, `x.t_()`, will change `x`.

```
1 print('x\n', x)
2 print('y\n', y)
```

```
x
tensor([[0.3215, 0.2230],
        [0.2855, 0.1580],
        [0.3406, 0.8168]])

y
tensor([[0.8634, 0.2672],
        [0.2046, 0.7790],
        [0.7036, 0.6863]])
```

```
1 y.add_(x)
2 print(y)
```

```
tensor([[1.1849, 0.4902],
        [0.4901, 0.9369],
        [1.0442, 1.5032]])
```

▼ Multiplication

Different from numpy which uses mainly `dot` to perform different types of matrix multiplication, PyTorch uses different commands for vector-vector multiplication (`dot`), matrix-vector multiplication (`mv`) and matrix-matrix multiplication (`mm`)

dot

```
1 a = torch.Tensor([4, 2])
2 b = torch.Tensor([3, 1])
3 r = torch.dot(a, b)
4
5 print(r)

tensor(14.)
```

mv

```
1 mat = torch.randn(2, 4)
2 vec = torch.randn(4)
3 r = torch.mv(mat, vec)
4 print(r)

tensor([ 2.8342, -1.1781])
```

mm

```
1 mat1 = torch.randn(2, 3)
2 mat2 = torch.randn(3, 4)
3 r = torch.mm(mat1, mat2)
4
5 print(r)

tensor([[ 0.6578, -1.6688, -0.2140,  0.7555],
        [-0.1618, -0.1385, -0.0776,  1.5601]])
```

▼ 3. Indexing

You can use standard Numpy-like indexing with Torch

```
1 x = torch.randint(0, 100, (5,10))
```

```
2 print(x)
```

```
tensor([[61, 31, 57, 48, 65, 98,  7, 13, 58, 14],
        [63,  7, 57, 32, 77, 44, 44, 71, 77, 32],
        [77, 60, 21, 68,  2, 13, 64, 74, 55, 33],
        [99, 74, 96, 71, 99, 25,  8, 77, 60, 70],
        [18, 31, 15,  6, 90, 12, 48, 81, 75, 62]])
```

```
1 # accessing column 1
```

```
2 print(x[:,1])
```

```
tensor([31,  7, 60, 74, 31])
```

```
1 # accessing columns 2 and 3
```

```
2 print(x[:, 2:4])
```

```
tensor([[57, 48],
        [57, 32],
        [21, 68],
        [96, 71],
        [15,  6]])
```

```
1 # accessing row 1
```

```
2 print(x[1,:])
```

```
tensor([63,  7, 57, 32, 77, 44, 44, 71, 77, 32])
```

```
1 # accessing rows 2 and 3
```

```
2 print(x[2:4,:])
```

```
tensor([[77, 60, 21, 68,  2, 13, 64, 74, 55, 33],
        [99, 74, 96, 71, 99, 25,  8, 77, 60, 70]])
```

```
tensor([[11, 88, 21, 88,  2, 13, 84, 74, 55, 55],
        [99, 74, 96, 71, 99, 25,  8, 77, 60, 70]])
```

▼ 4. Reshaping Tensors

Tensor.reshape

Returns a tensor with the same data and number of elements as self but with the specified shape.

```
1 x = torch.randint(0, 100, (2,4))
2 print('x:\n', x)
```

```
x:
tensor([[15, 97, 63, 37],
        [45, 41,  4, 97]])
```

```
1 # Reshape from (2, 4) to (8, 1)
2 y = x.reshape(8, -1)
3 print('y:\n', y)
```

```
y:
tensor([[15],
        [97],
        [63],
        [37],
        [45],
        [41],
        [ 4],
        [97]])
```

```
1 # Reshape from (2, 4) to (4, 2)
```

```
2 z = x.reshape(4, 2)
3 print('z:\n', z)
```

```
7.
```

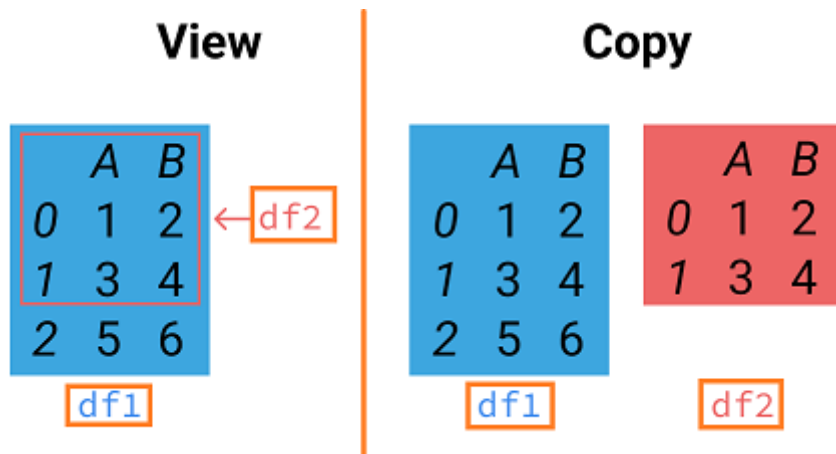


```

z =
  tensor([[15, 97],
         [63, 37],
         [45, 41],
         [ 4, 97]])

```

This method returns a **view** if shape is compatible with the current shape. Else, it may return a **copy**. This allows it to work with both [contiguous](#) and [non-contiguous](#) data.



In the examples above, we create a view since the shapes of reshaped tensors `y` and `z` are compatible with the original tensor `x`. Note that after a change is performed on `x`, then the changes will occur to both `y` and `z`.

The following code confirms that `y` and `z` are indeed **views** of `x`. Any changes to `x` will be observed in `y` and `z` as well.

```

1 x[0,0] = -3
2
3 print('x:\n', x)
4 print('y:\n', y)
5 print('z:\n', z)

x:
  tensor([[ -3, 97, 63, 37],
         [45, 41,  4, 97]])
y:
  tensor([[ -3],

```

```

    [97],
    [63],
    [37],
    [45],
    [41],
    [ 4],
    [97]])
z:
  tensor([[ -3, 97],
         [63, 37],
         [45, 41],
         [ 4, 97]])

```

Tensor.view

`Tensor.view` always returns a **view** of the original tensor with the new shape, i.e., it will share the underlying data with the original tensor.

```

1 x = torch.randint(0, 100, (2, 4))
2 print('x:\n', x)

```

```

x:
  tensor([[ 6, 94,  0, 65],
         [24, 59, 71, 69]])

```

```

1 # Convert from (2, 4) to (8, 1)
2 y = x.view(8, -1)
3 print('y:\n', y)
4

```

```

y:
  tensor([[ 6],
         [94],
         [ 0],

         [65],
         [24],
         [59],
         [71]])

```

```
[69]])
```

```
1 # Convert from (2, 4) to (4, 2)
2 z = x.view(4, 2)
3 print('z:\n', z)
```

```
z:
tensor([[ 6, 94],
        [ 0, 65],
        [24, 59],
        [71, 69]])
```

Similar to the numpy's `reshape` function, pytorch's `view` returns a reference of the original matrix albeit in a different shape

▼ 5. CUDA Tensors

Creating tensor in the GPU

```
1 if torch.cuda.is_available():
2     gpu = torch.device("cuda") # define a cuda device
3     x = torch.ones((2, 4), device = gpu) # Create the tensor in the GPU
4     print(x)

tensor([[1., 1., 1., 1.],
        [1., 1., 1., 1.]], device='cuda:0')
```

▼ Creating tensor in the cpu explicitly (default)

```
1 cpu = torch.device("cpu")
2 x = torch.ones((2, 4), device = cpu) # create a tensor in the CPU
3 print(x)
```

```
tensor([[1., 1., 1., 1.],
        [1., 1., 1., 1.]])
```

▼ Transferring tensor from cpu to gpu

Transfer using the `.cuda()` command.

```
1 x = torch.rand(3, 2, device = "cpu") # create tensor in cpu. The device argument also accepts a string besides a device object
2 print(x)
3
4 x = x.cuda() # move to GPU
5 print(x)

tensor([[0.6650, 0.4313],
        [0.0991, 0.8726],
        [0.5617, 0.5138]])
tensor([[0.6650, 0.4313],
        [0.0991, 0.8726],
        [0.5617, 0.5138]], device='cuda:0')
```

Transfer using the `.to()` command

```
1 x = torch.rand(3, 2) # create tensor in the CPU (default)
2 print(x)
3
4 gpu = torch.device('cuda') # move to GPU
5 x = x.to(gpu)
6 print(x)

tensor([[0.0312, 0.4889],
        [0.3286, 0.7417],
        [0.9349, 0.8032]])
tensor([[0.0312, 0.4889],
        [0.3286, 0.7417],
        [0.9349, 0.8032]], device='cuda:0')
```

▼ Transferring tensor from gpu to cpu

Transfer using the `.cpu()` command.

```
1 x = torch.rand(3, 2, device = 'cuda') # create tensor in gpu. The device argument also accepts a string besides a device object
2 print(x)
3
4 x = x.cpu() # move to CPU
5 print(x)

tensor([[0.9459, 0.8531],
        [0.0874, 0.3044],
        [0.0451, 0.5759]], device='cuda:0')
tensor([[0.9459, 0.8531],
        [0.0874, 0.3044],
        [0.0451, 0.5759]])
```

Transfer using the `.to()` command.

```
1 x = torch.rand(3, 2, device = 'cuda') # create tensor in the GPU
2 print(x)
3
4 device = torch.device('cpu')
5 x = x.to(device) # move to CPU
6 print(x)

tensor([[0.6613, 0.8984],
        [0.9873, 0.5311],
        [0.3895, 0.8701]], device='cuda:0')
tensor([[0.6613, 0.8984],
        [0.9873, 0.5311],
        [0.3895, 0.8701]])
```

▼ Exercise

Question 1. The following code is used to preprocess a batch data for Logistic Regression.

1.1 Create a random tensor `x_ori` using the normal distribution of shape `(4, 16, 16, 3)`. The tensor represent `m=4` color image samples, each having a resolution of `(16, 16)` Expected ans: Shape of `x_ori`: `torch.Size([4, 16, 16, 3])`

```
1 ...
2 print('Shape of X_ori:', X_ori.shape)
```

1.2 Reshape `x_ori` into a shape of `(4, 16*16*3)`. Then transpose the result to get a tensor of shape `(768, 4)` where each column represents a sample. Save the result as `x`.

```
Expected ans:
...
Shape of X: torch.Size([768, 4])
...
```

```
1 ...
2 print('Shape of X:', X.shape)
```

1.3 Check if a GPU is available in the system. If yes, transfer the tensor `x` to the GPU. Then, verify if `X` has really been loaded into the GPU (`x.is_cuda`) and print out the device ID of the GPU (`x.get_device()`).

Expected ans:
` `

```
x is loaded to GPU: 0
```

```
...
```

```
1 if ...gpu is available...
2     ... load x to gpu ...
3
4 if ...x is successfully loaded into GPU:
5     print('X is loaded to GPU:', ... get the GPU ID...)
6 else:
7     print('X is loaded to CPU')
```

Question 2.

2.1 Create the tensor A. Ensure that the datatype for A is `float32`:

```
A = [[3, 2, 4, 6],
      [2, 4, 2, 2],
      [5, 1, 2, 1]]
```

```
1 ...
2 print(A)
```

2.2 Extract the 2nd row from A. (Expected ans: `tensor([2., 4., 2., 2.])`)

```
1 print(...)
```

2.3 Extract the 3rd column from A. (Expected ans: `tensor([4., 2., 2.])`)

```
1 print(...)
```

2.4 Write the code to extract the following sub-block (rows 1 to 2 and columns 1 to 2) from A.

```
tensor([[4., 2.],
        [1., 2.]])
```

```
1 print(...)
```

2.5 Compute the mean of all rows. You may need to convert the tensor to a double first. Note that in the computational graph, when you create

Expected ans:

```
tensor([3.7500, 2.5000, 2.2500], dtype=torch.float64)
```

```
1 print(...)
```

2.6 Repeat question 2.5, but this time retain the original dimensions such that the output has a shape of (3,1)

Expected ans:

```
tensor([[3.7500],
        [2.5000],
        [2.2500]], dtype=torch.float64)
```

```
1 print(...)
```